

Classification Model Comparison: NYC Shooting Incidents

DTSA 5509 Introduction to Machine Learning, University of Colorado, Boulder, April 2023

Required Packages

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import math
import seaborn as sns
import sklearn
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn import metrics
import time
```

Problem Description

This project will analyze the NYPD Shooting Incident Data data set in an attempt to classify shooting incidents as murders. Three different classification models will be built and the classification accuracy compared against various model parameters. The K-Nearest Neighbors (KNN), Random Forest (Forest) and Gradient Boosting (Boost) Classifiers from the sklearn package will be evaluated. The features of the data set will be converted from multi-value to binary and a feature importance analysis will be performed. A comparison of computation time will be included to assess the best balance between performance and time.

The dataset being used is the NYPD Shooting Incident Data (Historic). The source file is located at <https://catalog.data.gov/dataset/nypd-shooting-incident-data-historic>. The data contains information on every shooting incident that occurred in New York City from 1/1/2006 through 12/31/2021. The data includes features related to the location and time of incident, perpetrator and victim attributes and murder label. The data set is comprised of 25,596 observations (rows) and 19 features (columns). The murder label (SATISTICAL_MURDER_FLAG) indicates a shooting that resulted in the victim's death and is identified as a murder. The ability to classify incidents as murders is important because it provides insight into the environmental conditions associated with a murder. Understanding these conditions may aid in avoiding these situations and reducing the number of murder incidents.

GITHUB Repository: <https://github.com/jmskeet/DTSA-5509>

EDA

Import Data and High Level Analysis

```
In [2]: originalDataSet = pd.read_csv('data/NYPD_Shooting_Incident_Data__Historic_.csv')
originalDataSet.describe()

print("Shape: " + str(originalDataSet.shape))
print("Row Count: " + str(len(originalDataSet)))
print(originalDataSet.head(5))
print(originalDataSet.describe(percentiles = [], include='all'))
print(originalDataSet.dtypes)
```

Shape: (25596, 19)

Row Count: 25596

	INCIDENT_KEY	OCCUR_DATE	OCCUR_TIME	BORO	PRECINCT	JURISDICTION_CODE	\
0	236168668	11/11/2021	15:04:00	BROOKLYN	79	0.0	
1	231008085	07/16/2021	22:05:00	BROOKLYN	72	0.0	
2	230717903	07/11/2021	01:09:00	BROOKLYN	79	0.0	
3	237712309	12/11/2021	13:42:00	BROOKLYN	81	0.0	
4	224465521	02/16/2021	20:00:00	QUEENS	113	0.0	

	LOCATION_DESC	STATISTICAL_MURDER_FLAG	PERP_AGE_GROUP	PERP_SEX	\
0	NaN	False	NaN	NaN	
1	NaN	False	45-64	M	
2	NaN	False	<18	M	
3	NaN	False	NaN	NaN	
4	NaN	False	NaN	NaN	

	PERP_RACE	VIC_AGE_GROUP	VIC_SEX	VIC_RACE	\
0	NaN	18-24	M	BLACK	
1	ASIAN / PACIFIC ISLANDER	25-44	M	ASIAN / PACIFIC ISLANDER	
2	BLACK	25-44	M	BLACK	
3	NaN	25-44	M	BLACK	
4	NaN	25-44	M	BLACK	

	X_COORD_CD	Y_COORD_CD	Latitude	Longitude	\
0	996313.0	187499.0	40.681318	-73.956509	
1	981845.0	171118.0	40.636364	-74.008667	
2	996546.0	187436.0	40.681145	-73.955669	
3	1001139.0	192775.0	40.695792	-73.939096	
4	1050710.0	184826.0	40.673740	-73.760411	

	Lon_Lat
0	POINT (-73.95650899099996 40.68131820000008)
1	POINT (-74.00866668999998 40.63636384100005)
2	POINT (-73.95566903799994 40.68114495900005)
3	POINT (-73.939095905 40.69579171600003)
4	POINT (-73.76041066999993 40.67374017600008)

	INCIDENT_KEY	OCCUR_DATE	OCCUR_TIME	BORO	PRECINCT	\
count	2.559600e+04	25596	25596	25596	25596.000000	
unique	NaN	5409	1411	5	NaN	
top	NaN	07/05/2020	23:30:00	BROOKLYN	NaN	
freq	NaN	47	171	10365	NaN	
mean	1.123826e+08	NaN	NaN	NaN	65.869433	
std	6.786117e+07	NaN	NaN	NaN	27.201904	
min	9.953245e+06	NaN	NaN	NaN	1.000000	
50%	8.643726e+07	NaN	NaN	NaN	69.000000	
max	2.384901e+08	NaN	NaN	NaN	123.000000	

	JURISDICTION_CODE	LOCATION_DESC	STATISTICAL_MURDER_FLAG	\
count	25594.000000	10619	25596	
unique	NaN	39	2	
top	NaN	MULTI DWELL - PUBLIC HOUS	False	
freq	NaN	4559	20668	
mean	0.331601	NaN	NaN	
std	0.742266	NaN	NaN	
min	0.000000	NaN	NaN	
50%	0.000000	NaN	NaN	
max	2.000000	NaN	NaN	

	PERP_AGE_GROUP	PERP_SEX	PERP_RACE	VIC_AGE_GROUP	VIC_SEX	VIC_RACE	\
count	16252	16286	16286	25596	25596	25596	

unique	9	3	7	6	3	7
top	18-24	M	BLACK	25-44	M	BLACK
freq	5844	14416	10668	11386	23182	18281
mean	NaN	NaN	NaN	NaN	NaN	NaN
std	NaN	NaN	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN	NaN	NaN

	X_COORD_CD	Y_COORD_CD	Latitude	Longitude	\
count	2.559600e+04	25596.000000	25596.000000	25596.000000	
unique	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	
mean	1.009455e+06	207893.776907	40.737250	-73.909039	
std	1.842142e+04	31857.353942	0.087447	0.066427	
min	9.149281e+05	125756.718750	40.511586	-74.249303	
50%	1.007715e+06	194037.718750	40.699128	-73.915346	
max	1.066815e+06	271127.687500	40.910818	-73.702046	

	Lon_Lat
count	25596
unique	11472
top	POINT (-73.88151014499994 40.67141260500006)
freq	66
mean	NaN
std	NaN
min	NaN
50%	NaN
max	NaN

INCIDENT_KEY	int64
OCCUR_DATE	object
OCCUR_TIME	object
BORO	object
PRECINCT	int64
JURISDICTION_CODE	float64
LOCATION_DESC	object
STATISTICAL_MURDER_FLAG	bool
PERP_AGE_GROUP	object
PERP_SEX	object
PERP_RACE	object
VIC_AGE_GROUP	object
VIC_SEX	object
VIC_RACE	object
X_COORD_CD	float64
Y_COORD_CD	float64
Latitude	float64
Longitude	float64
Lon_Lat	object
dtype:	object

High Level Observations

The following features contain NAs:

LOCATION_DESC
PERP_AGE_GROUP
PERP_SEX

PERP_RACE

The OCCUR_DATE and OCCUR_TIME are object types not datetime types.

The data will be assessed for the level of NA occurrences.

```
In [3]: nullValues = [np.nan, None, [], {}, 'NaN', 'Null', 'NULL', 'None', 'NA', '?', '-', '.', ''],
for c in originalDataSet.columns:
    string_null = np.array([x in nullValues[2:] for x in originalDataSet[c]])
    print(c, originalDataSet[c].isnull().sum(), string_null.sum())

INCIDENT_KEY 0 0
OCCUR_DATE 0 0
OCCUR_TIME 0 0
BORO 0 0
PRECINCT 0 0
JURISDICTION_CODE 2 0
LOCATION_DESC 14977 0
STATISTICAL_MURDER_FLAG 0 0
PERP_AGE_GROUP 9344 0
PERP_SEX 9310 0
PERP_RACE 9310 0
VIC_AGE_GROUP 0 0
VIC_SEX 0 0
VIC_RACE 0 0
X_COORD_CD 0 0
Y_COORD_CD 0 0
Latitude 0 0
Longitude 0 0
Lon_Lat 0 0
```

Based on the NA assessment the following features can be thrown. (NA count > 5%)

LOCATION_DESC
PERP_AGE_GROUP
PERP_SEX
PERP_RACE

Additionally, the following features are deemed either redundant or irrelevant to the analysis.

INCIDENT_KEY
PRECINCT
JURISDICTION_CODE
X_COORD_CD
Y_COORD_CD
Latitude
Longitude
Lon_Lat

```
In [4]: featuresToThrow = ['INCIDENT_KEY', 'PRECINCT', 'JURISDICTION_CODE', 'X_COORD_CD', 'Y_C
originalDataSet.drop(featuresToThrow, axis=1, inplace=True)
print("Shape: " + str(originalDataSet.shape))
```

```

print("Row Count: " + str(len(originalDataSet)))
print(originalDataSet.dtypes)

print ('')
print('NA Removal Validation')
for c in originalDataSet.columns:
    string_null = np.array([x in nullValues[2:] for x in originalDataSet[c]])
    print(c, originalDataSet[c].isnull().sum(), string_null.sum())

```

```

Shape: (25596, 7)
Row Count: 25596
OCCUR_DATE          object
OCCUR_TIME          object
BORO                object
STATISTICAL_MURDER_FLAG  bool
VIC_AGE_GROUP       object
VIC_SEX             object
VIC_RACE            object
dtype: object

```

```

NA Removal Validation
OCCUR_DATE 0 0
OCCUR_TIME 0 0
BORO 0 0
STATISTICAL_MURDER_FLAG 0 0
VIC_AGE_GROUP 0 0
VIC_SEX 0 0
VIC_RACE 0 0

```

The OCCUR_DATE is converted to date time object and broken into month and day features.
The OCCUR_TIME feature is broken out into hours of the day

```

In [5]: originalDataSet['OCCUR_DATE'] = pd.to_datetime(originalDataSet['OCCUR_DATE'])
originalDataSet['hour'] = originalDataSet['OCCUR_TIME'].astype(str).str[0]
originalDataSet['hour'] = originalDataSet['hour'].astype(int)
originalDataSet['day'] = originalDataSet['OCCUR_DATE'].dt.dayofweek
originalDataSet['month'] = originalDataSet['OCCUR_DATE'].dt.month
print(originalDataSet.dtypes)

```

```

OCCUR_DATE          datetime64[ns]
OCCUR_TIME          object
BORO                object
STATISTICAL_MURDER_FLAG  bool
VIC_AGE_GROUP       object
VIC_SEX             object
VIC_RACE            object
hour                int32
day                 int64
month               int64
dtype: object

```

Feature Value Assesment

Feature values will be evaluated via unique value and histogram assessment for completeness and correctness.

```

In [6]: boroughs = originalDataSet.BORO.unique()
murderFlag = originalDataSet.STATISTICAL_MURDER_FLAG.unique()

```


```
victimAgeGroup = originalDataSet.VIC_AGE_GROUP.unique()
victimSex = originalDataSet.VIC_SEX.unique()
victimRace = originalDataSet.VIC_RACE.unique()

print(boroughs)
print(murderFlag)
print(victimAgeGroup)
print(victimSex)
print(victimRace)

['BROOKLYN' 'QUEENS' 'BRONX' 'MANHATTAN' 'STATEN ISLAND']
[False True]
['18-24' '25-44' '<18' '45-64' '65+' 'UNKNOWN']
['M' 'F' 'U']
['BLACK' 'ASIAN / PACIFIC ISLANDER' 'BLACK HISPANIC' 'WHITE HISPANIC'
 'WHITE' 'AMERICAN INDIAN/ALASKAN NATIVE' 'UNKNOWN']
```


```
In [7]: plt.xticks(rotation='vertical')
ageHistogram = list(originalDataSet['VIC_AGE_GROUP'])
ageDict = {x:ageHistogram.count(x) for x in ageHistogram}
plt.bar(ageDict.keys(), ageDict.values(), color='g')
plt.legend(['Age'], loc=2)
```

```
Out[7]: <matplotlib.legend.Legend at 0x297c195ebc0>
```




```
In [8]: plt.xticks(rotation='vertical')
sexHistogram = list(originalDataSet['VIC_SEX'])
sexDict = {x:sexHistogram.count(x) for x in sexHistogram}
plt.bar(sexDict.keys(), sexDict.values(), color='b')
plt.legend(["Sex"], loc=1)
```

Out[8]: <matplotlib.legend.Legend at 0x297c1a572e0>



```
In [9]: plt.xticks(rotation='vertical')
raceHistogram = list(originalDataSet['VIC_RACE'])
raceDict = {x:raceHistogram.count(x) for x in raceHistogram}
plt.bar(raceDict.keys(), raceDict.values(), color='r')
plt.legend(["Race"], loc=1)
```

Out[9]: <matplotlib.legend.Legend at 0x297c195e7a0>



The feature value analysis showed that VIC_AGE_GROUP, VIC_SEX, VIC_RACE contained U and UNKNOWN Values. The degree to which these values impact the data set was assessed via percent of observations and the above histograms.

```
In [10]: unknownAgeRows = originalDataSet.apply(lambda x: True if x['VIC_AGE_GROUP'] == 'UNKNOWN' else False)
print('Unknown Age Rows: ' + str(len(unknownAgeRows[unknownAgeRows == True].index)))
print('Percent of Rows: ' + str(len(unknownAgeRows[unknownAgeRows == True].index)/len(originalDataSet)))

unknownSexRows = originalDataSet.apply(lambda x: True if x['VIC_SEX'] == 'U' else False)
print('Unknown Sex Rows: ' + str(len(unknownSexRows[unknownSexRows == True].index)))
print('Percent of Rows: ' + str(len(unknownSexRows[unknownSexRows == True].index)/len(originalDataSet)))

unknownRaceRows = originalDataSet.apply(lambda x: True if x['VIC_RACE'] == 'UNKNOWN' else False)
print('Unknown Race Rows: ' + str(len(unknownRaceRows[unknownRaceRows == True].index)))
print('Percent of Rows: ' + str(len(unknownRaceRows[unknownRaceRows == True].index)/len(originalDataSet)))
```

Unknown Age Rows: 60
Percent of Rows: 0.0023441162681669013
Unknown Sex Rows: 11
Percent of Rows: 0.0004297546491639319
Unknown Age Rows: 65
Percent of Rows: 0.0025394592905141427

The results revealed that <.5% of the observations are impacted and can be removed. Removal was validated.

```
In [11]: indexAge = originalDataSet[ (originalDataSet['VIC_AGE_GROUP'] == 'UNKNOWN')].index
originalDataSet.drop(indexAge, inplace=True)

indexSex = originalDataSet[ (originalDataSet['VIC_SEX'] == 'U')].index
originalDataSet.drop(indexSex, inplace=True)

indexRace = originalDataSet[ (originalDataSet['VIC_RACE'] == 'UNKNOWN')].index
originalDataSet.drop(indexRace, inplace=True)

#Verfiy rows removed
victimAgeGroup = originalDataSet.VIC_AGE_GROUP.unique()
victimSex = originalDataSet.VIC_SEX.unique()
victimRace = originalDataSet.VIC_RACE.unique()

print(victimAgeGroup)
print(victimSex)
print(victimRace)

print("Row Count: " + str(len(originalDataSet)))

['18-24' '25-44' '<18' '45-64' '65+']
['M' 'F']
['BLACK' 'ASIAN / PACIFIC ISLANDER' 'BLACK HISPANIC' 'WHITE HISPANIC'
 'WHITE' 'AMERICAN INDIAN/ALASKAN NATIVE']
Row Count: 25482
```

Data Cleansing

In summary, the data cleansing process consisted of identifying features and observations with NA values. Three features were identified as having excessive NAs (>5%) and were removed. Another nine features were identified and being redundant or irrelevaant to the analysis and also removed. The content of the remaining data set was then analyzed for content and correctness via unique value assessment and histograms. The content analysis identified three features with erroneous values. The number of observations related to these values totaled less than .5% of the observations and were removed.

The last step of the data analysis process was to convert the remaining features into a binary feature set. The resultant data set used to train and test the models consisted of 15,482 observations (rows) and 61 features (columns).

Build Classification Dataframe

With the original data set cleansed, the dataset was converted from multi-value features to binary features. The result being an increase of features from 10 to 61. For illustration purposes

the first 10 rows of the classification data set are provided.

```
In [12]: classificationData = pd.DataFrame()

classificationData['murder'] = originalDataSet['STATISTICAL_MURDER_FLAG']
#0 for female, 1 for male
classificationData['sex'] = [False if x == 'F' else True for x in originalDataSet['day']

#Map multivalue features to boolean feature columns
#day of week
for d in range(7):
    header = 'd' + str(d)
    classificationData[header] = [True if x == d else False for x in originalDataSet['c

#hour of day
for h in range(24):
    header = 'h' + str(h)
    classificationData[header] = [True if x == h else False for x in originalDataSet['h

#month of year
for m in range(1,13):
    header = 'm' + str(m)
    classificationData[header] = [True if x == m else False for x in originalDataSet['m

for b in boroughs:
    classificationData[b] = [True if x == b else False for x in originalDataSet['BORO']

for r in victimRace:
    classificationData[r] = [True if x == r else False for x in originalDataSet['VIC_RA

for a in victimAgeGroup :
    classificationData[a] = [True if x == a else False for x in originalDataSet['VIC_AC

classificationData.info()
print(classificationData.head(10))
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 25482 entries, 0 to 25595
```

```
Data columns (total 61 columns):
```

#	Column	Non-Null Count	Dtype
0	murder	25482 non-null	bool
1	sex	25482 non-null	bool
2	d0	25482 non-null	bool
3	d1	25482 non-null	bool
4	d2	25482 non-null	bool
5	d3	25482 non-null	bool
6	d4	25482 non-null	bool
7	d5	25482 non-null	bool
8	d6	25482 non-null	bool
9	h0	25482 non-null	bool
10	h1	25482 non-null	bool
11	h2	25482 non-null	bool
12	h3	25482 non-null	bool
13	h4	25482 non-null	bool
14	h5	25482 non-null	bool
15	h6	25482 non-null	bool
16	h7	25482 non-null	bool
17	h8	25482 non-null	bool
18	h9	25482 non-null	bool
19	h10	25482 non-null	bool
20	h11	25482 non-null	bool
21	h12	25482 non-null	bool
22	h13	25482 non-null	bool
23	h14	25482 non-null	bool
24	h15	25482 non-null	bool
25	h16	25482 non-null	bool
26	h17	25482 non-null	bool
27	h18	25482 non-null	bool
28	h19	25482 non-null	bool
29	h20	25482 non-null	bool
30	h21	25482 non-null	bool
31	h22	25482 non-null	bool
32	h23	25482 non-null	bool
33	m1	25482 non-null	bool
34	m2	25482 non-null	bool
35	m3	25482 non-null	bool
36	m4	25482 non-null	bool
37	m5	25482 non-null	bool
38	m6	25482 non-null	bool
39	m7	25482 non-null	bool
40	m8	25482 non-null	bool
41	m9	25482 non-null	bool
42	m10	25482 non-null	bool
43	m11	25482 non-null	bool
44	m12	25482 non-null	bool
45	BROOKLYN	25482 non-null	bool
46	QUEENS	25482 non-null	bool
47	BRONX	25482 non-null	bool
48	MANHATTAN	25482 non-null	bool
49	STATEN ISLAND	25482 non-null	bool
50	BLACK	25482 non-null	bool
51	ASIAN / PACIFIC ISLANDER	25482 non-null	bool
52	BLACK HISPANIC	25482 non-null	bool
53	WHITE HISPANIC	25482 non-null	bool
54	WHITE	25482 non-null	bool

```

55 AMERICAN INDIAN/ALASKAN NATIVE 25482 non-null bool
56 18-24 25482 non-null bool
57 25-44 25482 non-null bool
58 <18 25482 non-null bool
59 45-64 25482 non-null bool
60 65+ 25482 non-null bool
dtypes: bool(61)
memory usage: 1.7 MB
murder sex d0 d1 d2 d3 d4 d5 d6 h0 ... \
0 False True False False False True False False False False False ...
1 False True False False False False True False False False False ...
2 False True False False False False False False True True True ...
3 False True False False False False False True False False False ...
4 False True False True False False False False False False False ...
5 True True False False False False False True False False True ...
6 True True False False True False False False False False False ...
7 False True False False False False True False False False False ...
8 False True True False False False False False False False True ...
9 True True False False False False False False True True True ...

ASIAN / PACIFIC ISLANDER BLACK HISPANIC WHITE HISPANIC WHITE \
0 False False False False
1 True False False False
2 False False False False
3 False False False False
4 False False False False
5 False False False False
6 False False False False
7 False False False False
8 False True False False
9 False False True False

AMERICAN INDIAN/ALASKAN NATIVE 18-24 25-44 <18 45-64 65+
0 False True False False False False
1 False False True False False False
2 False False True False False False
3 False False True False False False
4 False False True False False False
5 False False True False False False
6 False True False False False False
7 False False True False False False
8 False False True False False False
9 False False True False False False

[10 rows x 61 columns]

```

Analysis

Models and Training

Train and test data sets were generated (75% train / 25% test) and the three models built with their respective outputs provided. The KNN model was assessed for 20 different k values beginning at 2. The Random Forest and Gradient Boost models were evaluated for 20 different learner values ranging from 25 to 500, at intervals of 50. Accuracy and computation times were collected for each model at each of the conditions.

Each models was evaluated at 20 different condition sets to facilitate a side by side comparison of classification accuracy and runtime.

The results of the evaluations are shown below.

```
In [13]: y = classificationData['murder']
X = classificationData.drop(['murder'], axis = 1)
X.info()
y.info()
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state=42)

dfAccuracy = pd.DataFrame()
dfTimes = pd.DataFrame()

#KNN Evaluation
knnStartTime = time.time()

knnScores = [None]
knnTimes =[None]
for k in range(2, 22):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    knnPredictions = knn.predict(X_test)
    knnScores.append(metrics.accuracy_score(y_test, knnPredictions))
    knnTimes.append(time.time() - knnStartTime)

dfAccuracy['KNN'] = knnScores
dfTimes['KNN'] = knnTimes

#Random Forest
forestStartTime = time.time()

estimators = [25, 50, 75, 100, 125, 150, 175, 200, 225, 250, 275, 300, 325, 350, 375,
forestScores = [None]
forestTimes = [None]
for e in estimators:
    forest = RandomForestClassifier(n_estimators=e)
    forest.fit(X_train, y_train)
    forestPredictions = forest.predict(X_test)
    forestScores.append(metrics.accuracy_score(y_test, forestPredictions))
    forestTimes.append(time.time() - forestStartTime)

dfAccuracy['Forest'] = forestScores
dfTimes['Forest'] = forestTimes

#Gradient Boost
boostStartTime = time.time()

boostScores = [None]
boostTimes = [None]
for e in estimators:
    boost = GradientBoostingClassifier(n_estimators=e)
    boost.fit(X_train, y_train)
    boostPredictions = boost.predict(X_test)
    boostScores.append(metrics.accuracy_score(y_test, boostPredictions))
    boostTimes.append(time.time() - boostStartTime)

dfAccuracy['Boost'] = boostScores
```

```
dfTimes['Boost'] = boostTimes

print("Table1: Accuracy Scores")
print(dfAccuracy)

print('')

print("Table2: Computation Times (seconds)")
print(dfTimes)
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 25482 entries, 0 to 25595
```

```
Data columns (total 60 columns):
```

#	Column	Non-Null Count	Dtype
----	-----	-----	-----
0	sex	25482 non-null	bool
1	d0	25482 non-null	bool
2	d1	25482 non-null	bool
3	d2	25482 non-null	bool
4	d3	25482 non-null	bool
5	d4	25482 non-null	bool
6	d5	25482 non-null	bool
7	d6	25482 non-null	bool
8	h0	25482 non-null	bool
9	h1	25482 non-null	bool
10	h2	25482 non-null	bool
11	h3	25482 non-null	bool
12	h4	25482 non-null	bool
13	h5	25482 non-null	bool
14	h6	25482 non-null	bool
15	h7	25482 non-null	bool
16	h8	25482 non-null	bool
17	h9	25482 non-null	bool
18	h10	25482 non-null	bool
19	h11	25482 non-null	bool
20	h12	25482 non-null	bool
21	h13	25482 non-null	bool
22	h14	25482 non-null	bool
23	h15	25482 non-null	bool
24	h16	25482 non-null	bool
25	h17	25482 non-null	bool
26	h18	25482 non-null	bool
27	h19	25482 non-null	bool
28	h20	25482 non-null	bool
29	h21	25482 non-null	bool
30	h22	25482 non-null	bool
31	h23	25482 non-null	bool
32	m1	25482 non-null	bool
33	m2	25482 non-null	bool
34	m3	25482 non-null	bool
35	m4	25482 non-null	bool
36	m5	25482 non-null	bool
37	m6	25482 non-null	bool
38	m7	25482 non-null	bool
39	m8	25482 non-null	bool
40	m9	25482 non-null	bool
41	m10	25482 non-null	bool
42	m11	25482 non-null	bool
43	m12	25482 non-null	bool
44	BROOKLYN	25482 non-null	bool
45	QUEENS	25482 non-null	bool
46	BRONX	25482 non-null	bool
47	MANHATTAN	25482 non-null	bool
48	STATEN ISLAND	25482 non-null	bool
49	BLACK	25482 non-null	bool
50	ASIAN / PACIFIC ISLANDER	25482 non-null	bool
51	BLACK HISPANIC	25482 non-null	bool
52	WHITE HISPANIC	25482 non-null	bool
53	WHITE	25482 non-null	bool
54	AMERICAN INDIAN/ALASKAN NATIVE	25482 non-null	bool

55	18-24	25482	non-null	bool
56	25-44	25482	non-null	bool
57	<18	25482	non-null	bool
58	45-64	25482	non-null	bool
59	65+	25482	non-null	bool

dtypes: bool(60)

memory usage: 1.7 MB

<class 'pandas.core.series.Series'>

Int64Index: 25482 entries, 0 to 25595

Series name: murder

Non-Null Count Dtype

25482 non-null bool

dtypes: bool(1)

memory usage: 224.0 KB

Table1: Accuracy Scores

	KNN	Forest	Boost
0	NaN	NaN	NaN
1	0.793753	0.778057	0.816198
2	0.753257	0.779783	0.815884
3	0.795636	0.781353	0.815414
4	0.777429	0.781196	0.815571
5	0.806938	0.784021	0.815728
6	0.798619	0.781039	0.815414
7	0.809606	0.783080	0.815414
8	0.805211	0.783707	0.815100
9	0.812431	0.782452	0.814943
10	0.807879	0.781196	0.814943
11	0.814315	0.781510	0.814472
12	0.812117	0.783237	0.815414
13	0.814158	0.783864	0.814629
14	0.813216	0.780097	0.814315
15	0.814786	0.783237	0.814943
16	0.813844	0.782923	0.814629
17	0.815414	0.784021	0.814472
18	0.814786	0.783707	0.814629
19	0.815414	0.785748	0.814315
20	0.815257	0.783550	0.814786

Table2: Computation Times (seconds)

	KNN	Forest	Boost
0	NaN	NaN	NaN
1	9.459464	0.513104	0.370999
2	13.943225	1.439506	1.057259
3	17.817199	2.843148	2.065574
4	20.152096	4.664503	3.407861
5	22.534063	6.580302	5.084746
6	24.904716	8.956427	7.086252
7	27.216801	11.630013	9.420906
8	29.862215	14.660146	12.265265
9	32.430725	18.452906	15.673562
10	34.998238	23.146643	19.097421
11	37.382952	27.688548	22.759451
12	39.727803	32.226657	26.769495
13	42.157475	38.155208	31.660589
14	44.579282	43.578819	36.462616
15	47.264323	49.512454	41.570634
16	49.893349	56.068776	47.704593
17	52.290138	62.400542	53.748124
18	54.608831	69.982752	60.023950

```
19 56.978517 77.036224 66.845310
20 59.330110 85.383639 73.572165
```


```
In [ ]: Plots were generated to compare models.
```

```
In [14]: xLabels = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20]
```


```
print('Plot 1: Accuracy Plot')
plt.plot(dfAccuracy['KNN'])
plt.plot(dfAccuracy['Forest'])
plt.plot(dfAccuracy['Boost'])
plt.xlabel('Varied Parameter Condition')
plt.ylabel('Accuracy')
plt.legend(['KNN', 'Forest', 'Boost'], loc=4)
plt.xticks(xLabels)
plt.show()
```

```
print('Plot 2: Computation Time Plot')
plt.plot(dfTimes['KNN'])
plt.plot(dfTimes['Forest'])
plt.plot(dfTimes['Boost'])
plt.xlabel('Condition Instance')
plt.ylabel('Seconds')
plt.legend(['KNN', 'Forest', 'Boost'], loc=4)
plt.xticks(xLabels)
plt.show()
```

Plot 1: Accuracy Plot



Plot 2: Computation Time Plot

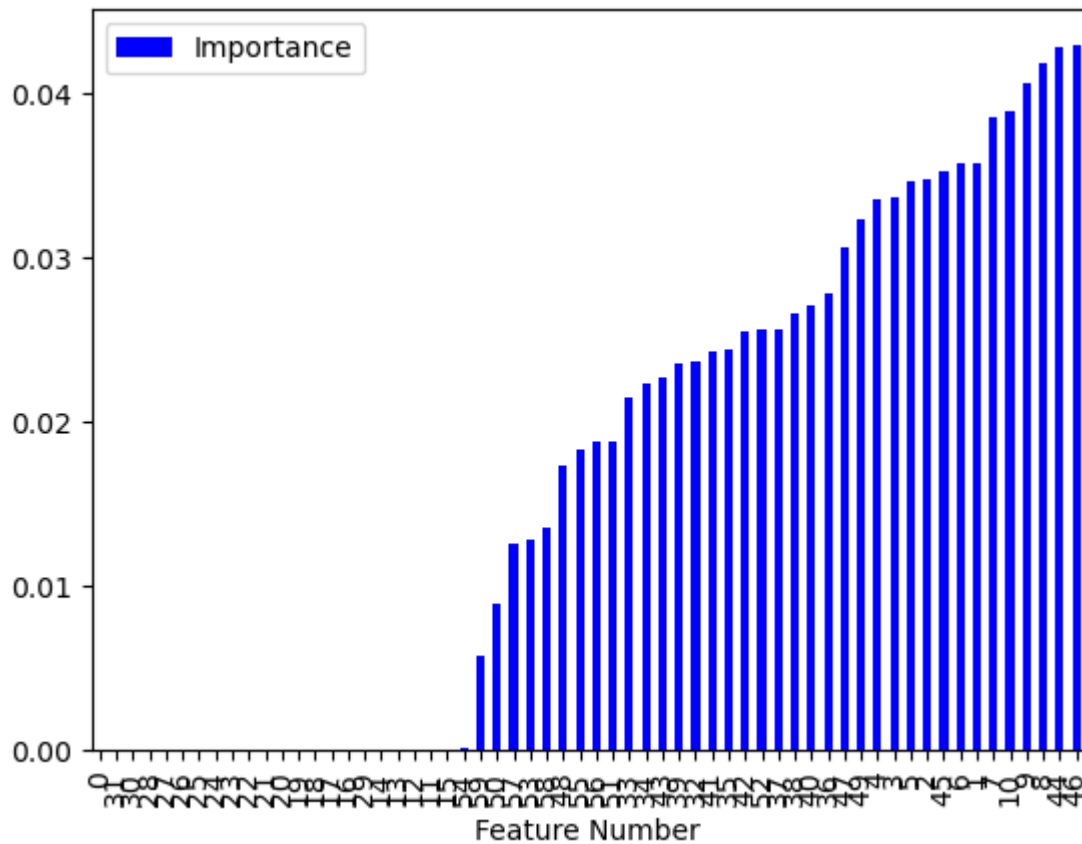


Feature Importance

A features importance test was performed on the dataset to assess the relevance of the 61 features and to provide insight into model performance.

```
In [15]: featureImportanceForest = RandomForestClassifier(n_estimators=300)
featureImportanceForest.fit(X_train, y_train)
featureImportance = featureImportanceForest.feature_importances_
importanceDf = pd.DataFrame({"Feature": pd.DataFrame(X_train).columns, "Importance":
importanceDf.set_index('Importance')
importanceDf = importanceDf.sort_values('Importance')
importanceDf.plot.bar(color = 'blue')
plt.xlabel('Feature Number')
plt.rc('font', size=6)
plt.show
```

```
Out[15]: <function matplotlib.pyplot.show(close=None, block=None)>
```



The feature importance assessment showed that sex and most of the hours of the day have no impact on the model. These features will be removed and the models run with the reduced feature set. For illustration purposes the first 10 rows of the classification data set are provided.

```
In [16]: featuresZeroImportance = importanceDf[importanceDf['Importance'] < .0001]
print(featuresZeroImportance)
featuresToRemove = list(featuresZeroImportance['Feature'])
print(featuresToRemove)

print('Create Important Feature Dataframe')
dfImportantFeatures = classificationData.drop(featuresToRemove, axis = 1)
dfImportantFeatures.info()
print(dfImportantFeatures.head(10))
```

	Feature	Importance
0	sex	0.0
31	h23	0.0
30	h22	0.0
28	h20	0.0
27	h19	0.0
26	h18	0.0
25	h17	0.0
24	h16	0.0
23	h15	0.0
22	h14	0.0
21	h13	0.0
20	h12	0.0
19	h11	0.0
18	h10	0.0
17	h9	0.0
16	h8	0.0
29	h21	0.0
14	h6	0.0
13	h5	0.0
12	h4	0.0
11	h3	0.0
15	h7	0.0

['sex', 'h23', 'h22', 'h20', 'h19', 'h18', 'h17', 'h16', 'h15', 'h14', 'h13', 'h12', 'h11', 'h10', 'h9', 'h8', 'h21', 'h6', 'h5', 'h4', 'h3', 'h7']

Create Important Feature Dataframe

<class 'pandas.core.frame.DataFrame'>

Int64Index: 25482 entries, 0 to 25595

Data columns (total 39 columns):

#	Column	Non-Null Count	Dtype
0	murder	25482 non-null	bool
1	d0	25482 non-null	bool
2	d1	25482 non-null	bool
3	d2	25482 non-null	bool
4	d3	25482 non-null	bool
5	d4	25482 non-null	bool
6	d5	25482 non-null	bool
7	d6	25482 non-null	bool
8	h0	25482 non-null	bool
9	h1	25482 non-null	bool
10	h2	25482 non-null	bool
11	m1	25482 non-null	bool
12	m2	25482 non-null	bool
13	m3	25482 non-null	bool
14	m4	25482 non-null	bool
15	m5	25482 non-null	bool
16	m6	25482 non-null	bool
17	m7	25482 non-null	bool
18	m8	25482 non-null	bool
19	m9	25482 non-null	bool
20	m10	25482 non-null	bool
21	m11	25482 non-null	bool
22	m12	25482 non-null	bool
23	BROOKLYN	25482 non-null	bool
24	QUEENS	25482 non-null	bool
25	BRONX	25482 non-null	bool
26	MANHATTAN	25482 non-null	bool
27	STATEN ISLAND	25482 non-null	bool
28	BLACK	25482 non-null	bool

29	ASIAN / PACIFIC ISLANDER	25482	non-null	bool
30	BLACK HISPANIC	25482	non-null	bool
31	WHITE HISPANIC	25482	non-null	bool
32	WHITE	25482	non-null	bool
33	AMERICAN INDIAN/ALASKAN NATIVE	25482	non-null	bool
34	18-24	25482	non-null	bool
35	25-44	25482	non-null	bool
36	<18	25482	non-null	bool
37	45-64	25482	non-null	bool
38	65+	25482	non-null	bool

dtypes: bool(39)

memory usage: 1.1 MB

	murder	d0	d1	d2	d3	d4	d5	d6	h0	h1	...	\
0	False	False	False	False	True	False	False	False	False	True	...	
1	False	False	False	False	False	True	False	False	False	False	...	
2	False	False	False	False	False	False	False	True	True	False	...	
3	False	False	False	False	False	False	True	False	False	True	...	
4	False	False	True	False	False	False	False	False	False	False	...	
5	True	False	False	False	False	False	True	False	True	False	...	
6	True	False	False	True	False	False	False	False	False	False	...	
7	False	False	False	False	False	True	False	False	False	True	...	
8	False	True	False	False	False	False	False	False	True	False	...	
9	True	False	False	False	False	False	False	True	True	False	...	

	ASIAN / PACIFIC ISLANDER	BLACK HISPANIC	WHITE HISPANIC	WHITE	\
0	False	False	False	False	
1	True	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	
5	False	False	False	False	
6	False	False	False	False	
7	False	False	False	False	
8	False	True	False	False	
9	False	False	True	False	

	AMERICAN INDIAN/ALASKAN NATIVE	18-24	25-44	<18	45-64	65+
0	False	True	False	False	False	False
1	False	False	True	False	False	False
2	False	False	True	False	False	False
3	False	False	True	False	False	False
4	False	False	True	False	False	False
5	False	False	True	False	False	False
6	False	True	False	False	False	False
7	False	False	True	False	False	False
8	False	False	True	False	False	False
9	False	False	True	False	False	False

[10 rows x 39 columns]

Reduced Feature Sets

Results of the models deployed with the reduced feature sets were compared.

```
In [17]: y = dfImportantFeatures['murder']
X = dfImportantFeatures.drop(['murder'], axis = 1)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state=42)

knnImportantFeaturesStartTime = time.time()
```

```

knnImportantFeaturesScores = [None]
knnImportantFeaturesTimes = [None]
for k in range(2, 22):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    knnImportantFeaturesPredictions = knn.predict(X_test)
    knnImportantFeaturesScores.append(metrics.accuracy_score(y_test, knnImportantFeaturesPredictions))
    knnImportantFeaturesTimes.append(time.time() - knnImportantFeaturesStartTime)

dfAccuracy['IF_KNN'] = knnImportantFeaturesScores
dfTimes['IF_KNN'] = knnImportantFeaturesTimes

forestImportantFeaturesStartTime = time.time()

forestImportantFeaturesScores = [None]
forestImportantFeaturesTimes = [None]
for e in estimators:
    forest = RandomForestClassifier(n_estimators=e)
    forest.fit(X_train, y_train)
    forestImportantFeaturesPredictions = forest.predict(X_test)
    forestImportantFeaturesScores.append(metrics.accuracy_score(y_test, forestImportantFeaturesPredictions))
    forestImportantFeaturesTimes.append(time.time() - forestImportantFeaturesStartTime)

dfAccuracy['IF_Forest'] = forestImportantFeaturesScores
dfTimes['IF_Forest'] = forestImportantFeaturesTimes

boostImportantFeaturesStartTime = time.time()

boostImportantFeaturesScores = [None]
boostImportantFeaturesTimes = [None]
for e in estimators:
    boost = GradientBoostingClassifier(n_estimators=e)
    boost.fit(X_train, y_train)
    boostImportantFeaturesPredictions = boost.predict(X_test)
    boostImportantFeaturesScores.append(metrics.accuracy_score(y_test, boostImportantFeaturesPredictions))
    boostImportantFeaturesTimes.append(time.time() - boostImportantFeaturesStartTime)

dfAccuracy['IF_Boost'] = boostImportantFeaturesScores
dfTimes['IF_Boost'] = boostImportantFeaturesTimes

print("Table 3: Accuracy Scores")
print(dfAccuracy)
print("")
print("Table 4: Computation Times (seconds)")
print(dfTimes)

```

Table 3: Accuracy Scores

	KNN	Forest	Boost	IF_KNN	IF_Forest	IF_Boost
0	NaN	NaN	NaN	NaN	NaN	NaN
1	0.793753	0.778057	0.816198	0.793753	0.774133	0.816041
2	0.753257	0.779783	0.815884	0.753257	0.777586	0.816041
3	0.795636	0.781353	0.815414	0.795636	0.779626	0.815414
4	0.777429	0.781196	0.815571	0.777429	0.779940	0.815728
5	0.806938	0.784021	0.815728	0.806938	0.780568	0.815571
6	0.798619	0.781039	0.815414	0.798619	0.781353	0.815414
7	0.809606	0.783080	0.815414	0.809606	0.785434	0.815257
8	0.805211	0.783707	0.815100	0.805211	0.781196	0.815100
9	0.812431	0.782452	0.814943	0.812431	0.781667	0.814472
10	0.807879	0.781196	0.814943	0.807879	0.781196	0.814472
11	0.814315	0.781510	0.814472	0.814315	0.785120	0.814315
12	0.812117	0.783237	0.815414	0.812117	0.782295	0.814315
13	0.814158	0.783864	0.814629	0.814158	0.783394	0.814315
14	0.813216	0.780097	0.814315	0.813216	0.782138	0.814315
15	0.814786	0.783237	0.814943	0.814786	0.784021	0.814943
16	0.813844	0.782923	0.814629	0.813844	0.781981	0.814315
17	0.815414	0.784021	0.814472	0.815414	0.785434	0.815100
18	0.814786	0.783707	0.814629	0.814786	0.783550	0.815100
19	0.815414	0.785748	0.814315	0.815414	0.784178	0.815100
20	0.815257	0.783550	0.814786	0.815257	0.783237	0.815571

Table 4: Computation Times (seconds)

	KNN	Forest	Boost	IF_KNN	IF_Forest	IF_Boost
0	NaN	NaN	NaN	NaN	NaN	NaN
1	9.459464	0.513104	0.370999	2.243302	0.555812	0.383451
2	13.943225	1.439506	1.057259	4.236195	1.436801	1.043539
3	17.817199	2.843148	2.065574	6.637040	2.596497	1.987026
4	20.152096	4.664503	3.407861	8.984363	4.052532	3.239850
5	22.534063	6.580302	5.084746	11.342106	5.886881	4.802998
6	24.904716	8.956427	7.086252	13.790931	8.104019	6.680939
7	27.216801	11.630013	9.420906	16.470638	10.659380	8.870945
8	29.862215	14.660146	12.265265	19.106115	13.781244	11.351934
9	32.430725	18.452906	15.673562	21.404692	17.598643	14.517940
10	34.998238	23.146643	19.097421	23.985020	21.286171	17.919112
11	37.382952	27.688548	22.759451	26.270485	25.288165	21.378161
12	39.727803	32.226657	26.769495	28.601827	29.880077	25.131454
13	42.157475	38.155208	31.660589	31.195630	35.003553	29.322651
14	44.579282	43.578819	36.462616	33.830776	40.066841	34.171435
15	47.264323	49.512454	41.570634	36.191046	45.648550	38.872572
16	49.893349	56.068776	47.704593	38.478483	51.939124	43.886303
17	52.290138	62.400542	53.748124	40.821111	58.072115	49.769670
18	54.608831	69.982752	60.023950	43.186188	65.207370	55.499724
19	56.978517	77.036224	66.845310	45.590299	72.133681	61.640319
20	59.330110	85.383639	73.572165	48.187767	80.032311	68.333158

Reduced Feature Plots

```
In [18]: plt.rc('font', size=8)

print('Plot 3: Reduced Feature Set Accuracy Plot')
plt.plot(dfAccuracy['IF_KNN'])
plt.plot(dfAccuracy['IF_Forest'])
plt.plot(dfAccuracy['IF_Boost'])
plt.xlabel('Condition Instance')
plt.ylabel('Accuracy')
plt.legend(['KNN', 'Forest', 'Boost'], loc=4)
```




```

plt.xticks(xLabels)
plt.show()


print('Plot 4: Reduced Feature Set Computation Time Plot')
plt.plot(dfTimes['IF_KNN'])
plt.plot(dfTimes['IF_Forest'])
plt.plot(dfTimes['IF_Boost'])
plt.xlabel('Condition Instance')
plt.ylabel('Seconds')
plt.legend(['KNN', 'Forest', 'Boost'], loc=4)
plt.xticks(xLabels)
plt.show()

```

Plot 3: Reduced Feature Set Accuracy Plot



Plot 4: Reduced Feature Set Computation Time Plot




Plots were generated for each model comparing performance between the data sets

KNN Plots


```
In [22]: print('Plot 5: KNN Comparison Accuracy Plot')
plt.plot(dfAccuracy['IF_KNN'], linestyle='dashed')
plt.plot(dfAccuracy['KNN'], linestyle='dotted')
plt.xlabel('Varied Parameter Condition')
plt.ylabel('Accuracy')
plt.legend(['Reduced Feature KNN', 'KNN'], loc=4)
plt.xticks(xLabels)
plt.show()

print('Plot 6: Reduced Feature Set Computation Time Plot')
plt.plot(dfTimes['IF_KNN'])
plt.plot(dfTimes['KNN'])
plt.xlabel('Varied Parameter Condition')
plt.ylabel('Seconds')
plt.legend(['Reduced Feature KNN', 'KNN'], loc=4)
plt.xticks(xLabels)
plt.show()
```

Plot 5: KNN Comparison Accuracy Plot



Plot 6: Reduced Feature Set Computation Time Plot



Random Forest Plots

```
In [23]: print('Plot 7: Reduced Feature Set Accuracy Plot')
plt.plot(dfAccuracy['IF_Forest'])
```


```

plt.plot(dfAccuracy['Forest'])
plt.xlabel('Varied Parameter Condition')
plt.ylabel('Accuracy')
plt.legend(['Reduced Feature Forest', 'Forest'], loc=4)
plt.xticks(xLabels)
plt.show()


print('Plot 8: Reduced Feature Set Computation Time Plot')
plt.plot(dfTimes['IF_Forest'])
plt.plot(dfTimes['Forest'])
plt.xlabel('Varied Parameter Condition')
plt.ylabel('Seconds')
plt.legend(['Reduced Feature Forest', 'Forest'], loc=4)
plt.xticks(xLabels)
plt.show()

```

Plot 7: Reduced Feature Set Accuracy Plot



Plot 8: Reduced Feature Set Computation Time Plot




Gradient Decent Plots


```
In [24]: print('Plot 9: Reduced Feature Set Accuracy Plot')
plt.plot(dfAccuracy['IF_Boost'])
plt.plot(dfAccuracy['Boost'])
plt.xlabel('Varied Parameter Condition')
plt.ylabel('Accuracy')
plt.legend(['Reduced Feature Boost', 'Boost'], loc=1)
plt.xticks(xLabels)
plt.show()

print('Plot 10: Reduced Feature Set Computation Time Plot')
plt.plot(dfTimes['IF_Boost'])
plt.plot(dfTimes['Boost'])
plt.xlabel('Varied Parameter Condition')
plt.ylabel('Seconds')
plt.legend(['Reduced Feature Boost', 'Boost'], loc=4)
plt.xticks(xLabels)
plt.show()
```

Plot 9: Reduced Feature Set Accuracy Plot



Plot 10: Reduced Feature Set Computation Time Plot



Discussion

Accuracy

Table 1 and Plot 1 provide a comparison of model accuracy based on the original data set of 61 features. The plot shows that the KNN and Boost models performed similarly. Both models resulted in an ultimate classification accuracy of 81%. The difference between the two models being that the Gradient Boost arrived at the 81% accuracy much more quickly than the KNN model. The Random Forest model performed the worst out of the three with its accuracy topping out at ~78%. Similar to the Gradient Boost model, the Random Forest the accuracy levels were arrived at more quickly than the KNN model.

The feature importance method from the Random Forest classifier showed that approximately 22 of the 61 features had no impact on the models. The features removed were the victim sex and the vast majority of the hours of the day. This makes sense since most of the murders occurred in the midnight and early morning hours and male victims far outweighed female victims. Expectations were that the models would perform better with the feature reduction. The models were rebuilt and run on the reduced feature set.

Table 3 and Plots 3, 5, 7 and 9 detail the accuracy comparisons between models, as well as individual model comparisons between feature set sizes. Expectations were that the accuracy would improve given the reduced feature set, however this was not the case. For all three models the accuracy between the feature set sizes were virtually identical. The KNN model resulted in an exact overlay and the tree models had minimal variations. These variations can be attributed to the randomness in the training and test sets used in the individual tree for each model.

Although not initially expected, the tracking of the accuracies between feature set sizes makes sense. For the KNN model, the unimportant features would have never fallen in the nearest neighbor group from the offset. They would reside outside the group as they had no influence on the model.

As for the tree models, the unimportant features would have been pruned out very early in the process, again with no influence on the model, removing them had no impact on the accuracy. Furthermore, the fact that the unimportant features were removed from the tree models in early in the classification, that explains why the variance in accuracy in those models was less than the KNN model. The KNN model had to progress through all ks looking at all points, whereas the tree models eliminate unimportant features from the offset.

Computational Time

Table 2 and Plot 2 provide insight into the computational resources, in terms of computation time, required for each model based on the original 61 features. The KNN model performed worst in the lower condition (lower k) state than both the Random Forest and Gradient Boost Models. However, at the $k \sim 12$ and learner count of ~300 all three models required about the same computational resources. This was the crossover point where the KNN model began to outperform the tree models requiring less resources as the k value and number of learners increased.

Table 4 and plots 4, 6, 8 and 10 provide a view of the computational resources required across and within models based on the feature set sizes. The results of these comparisons were more in line with initial expectations. Plot 4 shows that the comparison between models at the reduced feature set size had very similar shapes. However, the computational resources, as expected, were reduced by 9%, 10% and 28% for the KNN, Random Forest and Gradient Boost models respectively.

Plots 6, 8 and 10 provide a side by side comparison for each of the models. These plots clearly show the reduction on computational resources required for the reduced feature set size. Additionally, the magnitude of the changes falls in line with expectations. The KNN model still had to calculate up to 20 nearest neighbors, just with a reduced feature set, still fairly expensive as k grows. On the other hand, the tree models benefitted more from the reduced feature set size because the pruning did not have to initially remove the unimportant features and was thus able to get to the optimal condition more quickly.

Enhancements

One question that arises is how the conversion from multi-variable features to binary features impacted the outcome of this study. The ideal model contains a minimal number of features required to arrive at optimal performance. In this case, the move to binary features was purely for educational and experimental purposes. It would be very interesting to compare the current results to results using the original multi-value feature set. Forward or backward stepwise regression could be employed to identify the optimal feature set and then compare classification metrics between the two different approaches.

Conclusion

In the end, this project provided both expected and unexpected results. Classification accuracy comparisons were within expectations between models. All three models arrived at accuracies right around 80%. What was interesting was that the tree models converged to their maximum accuracy much more quickly than the KNN model. This is not surprising given the pruning characteristics in the tree models versus the iterative nature of the KNN model.

The reduction in computational resources based on the reduced feature set also fell in line with expectations. Reduction in the number of features resulted in on average a 23% decrease in computational resources.

Unexpected outcomes based on initial assumptions were the impacts (or lack of impact) based on the removal of the unimportant features. Initial expectations were that removing unimportant features would potentially promote increased accuracy. This was not the case but easily explainable once the structure of the models was considered.

When comparing all three models based on accuracy and resource requirements combined, the Gradient Boost model provided the best performance. The Gradient Boost model reached the highest accuracy while requiring the least amount of resources to get there. The KNN and Random Forest models were comparable. The KNN model provided a bit more accuracy but

required more resources than the Random Forest. On the other hand, the Random Forest arrived at its optimal accuracy more quickly than the KNN model, but suffered from a slightly degraded accuracy level.